

# Deep Learning Homework

## Deep Learning Házi feladat

Airbus Ship Detection Challenge

Nagy Dániel (UU5SCQ)  
Szénássy Márton (HIYXIQ)  
Varga-Labóczki Vazul (H947XW)

January 2, 2026

### Abstract

This report presents our solution to the Airbus Ship Detection Challenge using a U-Net architecture for ship segmentation in satellite images. We discuss our initial attempts with Segment Anything and the challenges we faced, leading us to adopt the U-Net model. Our implementation includes various callbacks to enhance training, and we address challenges such as class imbalance through data augmentation. Overall, our solution demonstrates the effectiveness of deep learning techniques in addressing real-world problems in maritime surveillance.

### Kivonat

Ebben a dolgozatban bemutatjuk a megoldásunkat az Airbus Ship Detection Challenge kihívásra, amelyben U-Net architektúrát alkalmaztunk a műholdfelvételeken megjelenő hajók szegmentálására. Ismertetjük kezdeti kísérleteinket a Segment Anything modellel, valamint azokat a kihívásokat, amelyekkel szembesültünk, és amelyek arra készítettek minket, hogy a U-Net modellt válasszuk. Megvalósításunk különböző callback mechanizmusokat tartalmaz a tanulási folyamat javítása érdekében, és foglalkozunk olyan kihívásokkal, mint az osztályok közötti egyensúlyhiány, amelyet adatnövelési technikákkal kezelünk. Összességében megoldásunk bemutatja a mélytanulási technikák hatékonyságát a tengeri megfigyelés valós problémáinak kezelésében.

### Contents

24	<b>1 Introduction</b>	2
25	<b>2 Description of topic and previous solutions</b>	2
26	<b>3 Architecture</b>	2
27	3.1 First attempt . . . . .	2
28	3.2 Final model . . . . .	2
29	3.2.1 Model . . . . .	2
30	3.2.2 Callbacks . . . . .	3

31	<b>4 Implementation</b>	3
32	4.1 Training . . . . .	3
33	4.1.1 Dataset . . . . .	3
34	4.1.2 Preprocessing . . . . .	4
35	4.1.3 Problems . . . . .	4
36	<b>5 Summary</b>	4

## 37 **1 Introduction**

38 This report presents our solution to the Airbus Ship Detection Challenge [1]. This challenge focuses  
 39 on the task of detecting ships in satellite images using deep learning techniques [2, 7]. The challenge  
 40 addresses a real-world problem in maritime surveillance and ocean monitoring, where automated  
 41 detection systems can significantly reduce manual labor and improve response times. We solved this  
 42 problem by developing a convolutional neural network-based segmentation model that accurately  
 43 identifies ship locations in satellite imagery.

## 44 **2 Description of topic and previous solutions**

45 The Airbus Ship Detection Challenge is a competition hosted on the Kaggle platform, aiming to  
 46 develop algorithms for detecting ships in satellite images. The challenge provides a dataset of  
 47 high-resolution satellite images, each annotated with masks indicating the presence and location  
 48 of ships. Participants are tasked with creating models that can accurately segment ships from the  
 49 background in these images.  
 50 Previous solutions include various deep learning architectures, such as U-Net [8], Mask R-CNN [3]  
 51 and Segment Anything [6]. These models leverage convolutional neural networks (CNNs) to learn  
 52 spatial hierarchies of features from the input images. Many participants have also employed data  
 53 augmentation techniques to enhance model robustness and improve generalization to unseen data.

## 54 **3 Architecture**

### 55 **3.1 First attempt**

56 Our initial approach involved implementing Segment Anything [6], a state-of-the-art segmentation  
 57 model known for its versatility and performance across various segmentation tasks. However, we  
 58 encountered significant challenges during this phase due to the model’s way of working. Segment  
 59 Anything is designed to generate segmentation masks based on user-provided prompts, such as points  
 60 or bounding boxes. This interactive nature made it difficult to adapt the model for fully automated ship  
 61 detection in satellite images, as required by the challenge. This made it impractical for our specific  
 62 use case, leading us to explore alternative architectures better suited for automated segmentation  
 63 tasks.

### 64 **3.2 Final model**

#### 65 **3.2.1 Model**

66 After evaluating various architectures, we decided to implement a U-Net model [8] for our ship  
 67 detection task. U-Net is a convolutional neural network architecture specifically designed for  
 68 biomedical image segmentation but has proven effective in various segmentation tasks, including  
 69 satellite imagery analysis. The U-Net architecture consists of a contracting path (encoder) and an  
 70 expansive path (decoder). The encoder captures context and features from the input images through a

71 series of convolutional and pooling layers [9], while the decoder reconstructs the segmentation mask  
72 using upsampling and convolutional layers [10].

```
class SegmentationModel():
    def __init__(self):
        self.model = Sequential()
        # Encoder
        self.model.add(Conv2D(filters=16, kernel_size=3, activation='relu', padding = 'same', input_shape=(IMG_HEIGHT,IMG_WIDTH,3)))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=32, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=32, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=64, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=64, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=128, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=128, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=256, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=256, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        # Decoder
        self.model.add(Conv2DTranspose(filters=128, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=64, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=32, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=16, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        # Classifier
        self.model.add(Conv2D(filters=1, kernel_size=5, padding = 'same'))
```

Figure 1: Model architecture.

### 73 3.2.2 Callbacks

74 To enhance the training process and improve model performance, we incorporated several callbacks  
75 into our training routine:

- 76 • **Model Checkpointing:** We used model checkpointing to save the best model weights based  
77 on validation loss during training. This ensures that we retain the most effective model  
78 configuration.
- 79 • **Early Stopping:** Early stopping was implemented to monitor the validation loss and halt  
80 training if no improvement was observed for a specified number of epochs. This helps  
81 prevent overfitting and reduces unnecessary computation.
- 82 • **Learning Rate Reduction:** We employed a learning rate reduction strategy that decreases  
83 the learning rate when the validation loss plateaus. This allows the model to fine-tune its  
84 weights more effectively during later stages of training.

## 85 4 Implementation

### 86 4.1 Training

87 We trained our U-Net model using the Adam optimizer [5] with a learning rate of 0.003 and a batch  
88 size of 32. The training was conducted on our own hardware, utilizing a GPU to accelerate the  
89 training process. We employed the binary cross-entropy loss function, which is suitable for binary  
90 segmentation tasks like ship detection. The model was trained for 1000 epochs, with early stopping  
91 implemented to prevent overfitting.

#### 92 4.1.1 Dataset

93 We utilized the Airbus Ship Detection Challenge dataset [1], which comprises high-resolution satellite  
94 images annotated with ship masks. The dataset includes a diverse range of images, capturing various  
95 sea conditions, ship sizes, and orientations.

96 **4.1.2 Preprocessing**

97 To enhance the model's performance, we incorporated data augmentation techniques during training,  
98 such as random rotations, flips, and zooms. This is the main task of the preprocessing step, as it  
99 increases the diversity of the training data and helps prevent overfitting. Additionally, we applied  
100 batch normalization [4] to stabilize the training process.

101 **4.1.3 Problems**

102 During the training process, we encountered several challenges that required careful consideration  
103 and adjustments to our approach. One of the primary issues was dealing with class imbalance in the  
104 dataset, as the number of pixels representing ships was significantly lower than the background pixels.  
105 In the begining we solved this problem by providing more images with ships during training. In  
106 later trainings however, this problem was mitigated enough by data augmentation techniques, which  
107 helped to create a more balanced representation of ship pixels in the training data and the model was  
108 able to learn.

109 **5 Summary**

110 In this report, we presented our approach to the Airbus Ship Detection Challenge using a U-Net  
111 architecture for ship segmentation in satellite images. We discussed our initial attempts with Segment  
112 Anything and the challenges we faced, leading us to adopt the U-Net model. Our implementation  
113 included various callbacks to enhance training, and we addressed challenges such as class imbalance  
114 through data augmentation. Overall, our solution demonstrates the effectiveness of deep learning  
115 techniques in addressing real-world problems in maritime surveillance.

116 **References**

- 117 [1] Airbus. Airbus ship detection challenge. Kaggle, 2018. URL <https://www.kaggle.com/c/airbus-ship-detection>.
- 118
- 119 [2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- 120 [3] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn. In *IEEE International Conference on Computer Vision*, pages 2961–2969, 2017.
- 121
- 122 [4] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International Conference on Machine Learning*, pages 448–456, 2015.
- 123
- 124 [5] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- 125
- 126 [6] Alexander Kirillov, Eric Mintun, Nikhila Ravi, et al. Segment anything. *arXiv preprint arXiv:2304.02643*, 2023.
- 127
- 128 [7] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- 129
- 130 [8] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical Image Computing and Computer-Assisted Intervention*, pages 234–241, 2015.
- 131
- 132 [9] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- 133
- 134 [10] Jost Tobias Springenberg, Alexey Dosovitskiy, Thomas Brox, and Martin Riedmiller. Striving for simplicity: The all convolutional net. *arXiv preprint arXiv:1412.1535*, 2015.
- 135
- 136
- 137