

Deep Learning Homework

Deep Learning Házi feladat Airbus Ship Detection Challenge

Nagy Dániel (UU5SCQ)
Szénássy Márton (HIYXIQ)
Varga-Labóczki Vazul (H947XW)

January 7, 2026

Abstract

This report presents our solution to the Airbus Ship Detection Challenge using a fully convolutional encoder-decoder architecture for ship segmentation in satellite images. We discuss our initial attempts with Segment Anything and the challenges we faced, leading us to adopt the FCN model. Our implementation includes various callbacks to enhance training, and we address challenges such as class imbalance through data augmentation. Overall, our solution demonstrates the effectiveness of deep learning techniques in addressing real-world problems in maritime surveillance.

Kivonat

Ebben a dolgozatban bemutatjuk a megoldásunkat az Airbus Ship Detection Challenge kihívásra, amelyben FCN architektúrát alkalmaztunk a műholdfelvételeken megjelenő hajók szegmentálására. Ismertetjük kezdeti kísérleteinket a Segment Anything modellel, valamint azokat a kihívásokat, amelyekkel szembesültünk, és amelyek arra készítettek minket, hogy a FCN modellt válasszuk. Megvalósításunk különböző callback mechanizmusokat tartalmaz a tanulási folyamat javítása érdekében, és foglalkozunk olyan kihívásokkal, mint az osztályok közötti egyensúlyhiány, amelyet adat augmentációs technikákkal kezelünk. Összességében megoldásunk bemutatja a mélytanulási technikák hatékonysegétségét a tengeri megfigyelés valós problémáinak kezelésében.

Contents

25	1 Introduction	2
26	2 Description of topic and previous solutions	2
27	3 Architecture	2
28	3.1 First attempt	2
29	3.2 Final model	2
30	3.2.1 Callbacks	3

31	4 Implementation	4
32	4.1 Training	4
33	4.1.1 Preprocessing	4
34	4.1.2 Problems	4
35	5 Summary	4

36 **1 Introduction**

37 This report presents our solution to the Airbus Ship Detection Challenge [1]. This challenge focuses
 38 on the task of detecting ships in satellite images using deep learning techniques [3, 8]. The challenge
 39 addresses a real-world problem in maritime surveillance and ocean monitoring, where automated
 40 detection systems can significantly reduce manual labor and improve response times. We solved this
 41 problem by developing a convolutional neural network-based segmentation model that accurately
 42 identifies ship locations in satellite imagery, drawing on recent advances in fully convolutional
 43 networks [9].

44 **2 Description of topic and previous solutions**

45 The Airbus Ship Detection Challenge is a competition hosted on the Kaggle platform, aiming to
 46 develop algorithms for detecting ships in satellite images. The challenge provides a dataset of
 47 high-resolution satellite images, each annotated with masks indicating the presence and location
 48 of ships. Participants are tasked with creating models that can accurately segment ships from the
 49 background in these images.
 50 Previous solutions include various deep learning architectures, such as U-Net [11], Mask R-CNN [4],
 51 DeepLabv3+ [2], and Segment Anything [7]. These models leverage convolutional neural networks
 52 (CNNs) to learn spatial hierarchies of features from the input images, often building on strong
 53 backbones like VGG [13]. Many participants have also employed data augmentation techniques to
 54 enhance model robustness and improve generalization to unseen data [10, 12].

55 **3 Architecture**

56 **3.1 First attempt**

57 Our initial approach involved implementing Segment Anything [7], a state-of-the-art segmentation
 58 model known for its versatility and performance across various segmentation tasks. However, we
 59 encountered significant challenges during this phase due to the model’s way of working. Segment
 60 Anything is designed to generate segmentation masks based on user-provided prompts, such as points
 61 or bounding boxes. This interactive nature made it difficult to adapt the model for fully automated ship
 62 detection in satellite images, as required by the challenge. This made it impractical for our specific
 63 use case, leading us to explore alternative architectures better suited for automated segmentation
 64 tasks.

65 **3.2 Final model**

66 After evaluating various architectures, we implemented a fully convolutional network (FCN). The
 67 chosen FCN is a simple encoder–decoder network composed exclusively of convolutional layers
 68 (no dense layers), with learned downsampling via strided convolutions and learned upsampling
 69 via transpose convolutions. Unlike U-Net, this architecture does not use skip-connections between
 70 encoder and decoder paths.

- 71 Key characteristics of our FCN:
- 72 • Encoder: several Conv2D blocks with ReLU activations and BatchNormalization [5]. Down-
73 sampling is performed using strided convolutions that reduce spatial resolution (e.g. 256 ->
74 128 -> 64 -> 32 -> 16).
- 75 • Decoder: Conv2DTranspose layers with ReLU activations and BatchNormalization [5] to
76 restore spatial resolution back to the input size.
- 77 • Classifier: a final Conv2D producing a single-channel output (kernel size 5, padding='same').
78 The model emits logits for binary segmentation; a sigmoid can be applied at loss/metric
79 time.
- 80 • BatchNormalization is applied after each hidden convolutional layer to stabilize training.

```

class SegmentationModel():
    def __init__(self):
        self.model = Sequential()
        # Encoder
        self.model.add(Conv2D(filters=16, kernel_size=3, activation='relu', padding = 'same', input_shape=(IMG_HEIGHT,IMG_WIDTH,3)))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=32, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=32, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=64, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=64, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=128, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=128, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=256, kernel_size=3, activation='relu', padding = 'same', strides = 2))
        self.model.add(BatchNormalization())
        self.model.add(Conv2D(filters=256, kernel_size=3, activation='relu', padding = 'same'))
        self.model.add(BatchNormalization())

        # Decoder
        self.model.add(Conv2DTranspose(filters=128, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=64, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=32, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())
        self.model.add(Conv2DTranspose(filters=16, kernel_size=3, strides=2, padding='same', activation='relu'))
        self.model.add(BatchNormalization())

        # Classifier
        self.model.add(Conv2D(filters=1, kernel_size=5, padding = 'same'))

```

Figure 1: Model architecture (fully convolutional encoder-decoder).

- 81 **3.2.1 Callbacks**
- 82 To enhance the training process and improve model performance, we incorporated several callbacks
83 into our training routine:
- 84 • **Model Checkpointing:** We used model checkpointing to save the best model weights based
85 on validation loss during training. This ensures that we retain the most effective model
86 configuration.
- 87 • **Early Stopping:** Early stopping was implemented to monitor the validation loss and halt
88 training if no improvement was observed for a specified number of epochs. This helps
89 prevent overfitting and reduces unnecessary computation.
- 90 • **Learning Rate Reduction:** We employed a learning rate reduction strategy that decreases
91 the learning rate when the validation loss plateaus. This allows the model to fine-tune its
92 weights more effectively during later stages of training.

93 **4 Implementation**

94 **4.1 Training**

95 We trained our FCN model using the Adam optimizer [6] with a learning rate of 0.003 and a batch
96 size of 32. The training was conducted on our own hardware, utilizing a GPU to accelerate the
97 training process. The model was trained for 1000 epochs, with early stopping implemented to prevent
98 overfitting.

99 **4.1.1 Preprocessing**

100 To enhance the model’s performance, we incorporated data augmentation techniques during training,
101 such as random rotations, flips, and zooms [10, 12]. This is the main task of the preprocessing step,
102 as it increases the diversity of the training data and helps prevent overfitting. Additionally, we applied
103 batch normalization [5] to stabilize the training process.

104 **4.1.2 Problems**

105 During the training process, we encountered several challenges that required careful consideration
106 and adjustments to our approach. One of the primary issues was dealing with class imbalance in the
107 dataset, as the number of pictures containing ships was significantly lower than the pictures without
108 ships. To address this, we implemented data augmentation techniques to increase the representation
109 of ship-containing images in the training set. The final ratio is 50-50 between images with and
110 without ships.

111 **5 Summary**

112 In this report, we presented our approach to the Airbus Ship Detection Challenge using a fully convolutional
113 encoder-decoder architecture for ship segmentation in satellite images. We discussed our initial
114 attempts with Segment Anything and the challenges we faced, leading us to adopt the FCN model.
115 Our implementation included various callbacks to enhance training, and we addressed challenges such
116 as class imbalance through data augmentation. Overall, our solution demonstrates the effectiveness
117 of deep learning techniques in addressing real-world problems in maritime surveillance [9, 2, 14].

118 **References**

- 119 [1] Airbus. Airbus ship detection challenge. Kaggle, 2018. URL <https://www.kaggle.com/c/airbus-ship-detection>.
- 121 [2] Liang-Chieh Chen, Yukun Zhu, George Papandreou, Florian Schroff, and Hartwig Adam.
122 Encoder-decoder with atrous separable convolution for semantic image segmentation. In
123 *European Conference on Computer Vision*, pages 833–851, 2018.
- 124 [3] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- 125 [4] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn. In *IEEE
126 International Conference on Computer Vision*, pages 2961–2969, 2017.
- 127 [5] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training
128 by reducing internal covariate shift. In *International Conference on Machine Learning*, pages
129 448–456, 2015.
- 130 [6] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint
131 arXiv:1412.6980*, 2014.
- 132 [7] Alexander Kirillov, Eric Mintun, Nikhila Ravi, et al. Segment anything. *arXiv preprint
133 arXiv:2304.02643*, 2023.

- 134 [8] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning
135 applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- 136 [9] Jonathan Long, Evan Shelhamer, and Trevor Darrell. Fully convolutional networks for semantic
137 segmentation. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 3431–
138 3440, 2015.
- 139 [10] Luis Pérez and Jason Wang. The effectiveness of data augmentation in image classification
140 using deep learning. *arXiv preprint arXiv:1712.04621*, 2017.
- 141 [11] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for
142 biomedical image segmentation. In *International Conference on Medical Image Computing and
143 Computer-Assisted Intervention*, pages 234–241, 2015.
- 144 [12] Connor Shorten and Taghi M. Khoshgoftaar. A survey on image data augmentation for deep
145 learning. *Journal of Big Data*, 6(60), 2019.
- 146 [13] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale
147 image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- 148 [14] Zongwei Zhou, Md Mahfuzur Rahman Siddiquee, Nima Tajbakhsh, and Jianming Liang.
149 Unet++: A nested u-net architecture for medical image segmentation. In *Deep Learning in
150 Medical Image Analysis and Multimodal Learning for Clinical Decision Support*, pages 3–11,
151 2018.